

# HIERARCHICAL MODELS I

DR. OLANREWAJU MICHAEL AKANDE

FEB 28, 2020

# ANNOUNCEMENTS

- No HW today.
- Next HW immediately after spring break on Monday, March 16.
- Midterm exam next Friday, March 6.
- Practice questions on Sakai later today or tomorrow.
- Review session next Wednesday, March 4.

# OUTLINE

- Introduction to hierarchical models
- Shrinkage
- Comparing two groups
- BMI example
- Comparing multiple groups with same variance

# MOTIVATION

- Sometimes, we may have a natural grouping in our data, for example
  - students within schools,
  - patients within hospitals,
  - voters within counties or states,
  - biology data, where animals are followed within natural populations organized geographically and, in some cases, socially.
- For such grouped data, we may want to do inference across all the groups, for example, comparison of the group means.
- Ideally, we should do so in a way that takes advantage of the relationship between observations in the same group, but we should also look to borrow information across groups when possible.
- **Hierarchical modeling** provides a principled way to do so.

# BAYES ESTIMATORS AND BIAS

- Recall the normal model:

$$y_i | \mu, \sigma^2 \stackrel{iid}{\sim} \mathcal{N}(\mu, \sigma^2).$$

- The MLE for the population mean  $\mu$  is just the sample mean  $\bar{y}$ .
- $\bar{y}$  is unbiased for  $\mu$ . That is, for any data  $y_i \stackrel{iid}{\sim} \mathcal{N}(\mu, \sigma^2)$ ,  $\mathbb{E}[\bar{y}] = \mu$ .
- However, recall that in the conjugate Normal-Gamma normal for example, the posterior expectation is a **weighted average** of the prior mean and the sample mean.
- That is, it is actually biased!

# SHRINKAGE

- Usually through the weighting of the sample data and prior, the Bayes procedure has the tendency to pull the estimate of  $\mu$  toward the prior mean.
- Of course, the magnitude of the pull depends on the sample size.
- This "pulling" phenomenon is referred to as **shrinkage**.
- Why would we ever want to do this? Why not just stick with the MLE?
- Well, in part, because shrinkage estimators are often "more accurate" in prediction problems – i.e. they tend to do a better job of predicting a future outcome or of recovering the actual parameter values. Remember variance-bias trade off!
- The fact that a biased estimator would do a better job in many prediction problems can be proven rigorously, and is referred to as **Stein's paradox**.

# MODERN RELEVANCE

- Stein's result implies, in particular, that the sample mean is an *inadmissible* estimator of the mean of a multivariate normal distribution in more than two dimensions – i.e. there are other estimators that will come closer to the true value in expectation.
- In fact, these are Bayes point estimators (the posterior expectation of the parameter  $\mu$  ).
- Most of what we do now in high-dimensional statistics is develop biased estimators that perform better than unbiased ones.
- Examples: lasso regression, ridge regression, various kinds of hierarchical Bayesian models, etc.
- Today we will get a very basic introduction to **Bayesian hierarchical models**, which provide a formal and coherent framework for constructing shrinkage estimators.

# WHY HIERARCHICAL MODELS?

- **Bayesian hierarchical models** is a sort of catch-all phrase for a large class of models that have several levels of conditional distributions making up the prior.
- Like simpler one-level priors, they also accomplish shrinkage. However, they are much more flexible.
- Why use them? Several reasons:
  - We may want to exploit more complex dependence structures.
  - We may have many parameters relative to the amount of data that we have, and want to borrow information in estimating them.
  - We may want to shrink toward something other than a simple prior mean/hyper-parameter.



# COMPARING TWO GROUPS

- Suppose we want to do inference on mean body mass index (BMI) for two groups (male or female).
- BMI is known to often follow a normal distribution, so let's assume the same here.
- We should expect some relationship between the mean BMI for the two groups.
- We may also think the shape of the two distributions would be relatively the same (at least as a simplifying assumption for now).
- Thus, a reasonable model might be

$$\begin{aligned} y_{i,\text{male}} &\stackrel{iid}{\sim} \mathcal{N}(\theta_m, \sigma^2); \quad i = 1, \dots, n_m; \\ y_{i,\text{female}} &\stackrel{iid}{\sim} \mathcal{N}(\theta_f, \sigma^2); \quad i = 1, \dots, n_f. \end{aligned}$$

but with some relationship between  $\theta_m$  and  $\theta_f$ .

# APPLICATION

- First, let's do classical inference on such data. The data we will use in the R package `rethinking`.

```
#install.packages(c("coda", "mvtnorm", "devtools", "loo", "dagitty"))
#library(devtools)
#devtools::install_github("rmcelreath/rethinking", ref="Experimental")
#library(rethinking)
data(Howell1)

Howell1[1:15,]
```

##		height	weight	age	male
## 1		151.765	47.82561	63.0	1
## 2		139.700	36.48581	63.0	0
## 3		136.525	31.86484	65.0	0
## 4		156.845	53.04191	41.0	1
## 5		145.415	41.27687	51.0	0
## 6		163.830	62.99259	35.0	1
## 7		149.225	38.24348	32.0	0
## 8		168.910	55.47997	27.0	1
## 9		147.955	34.86988	19.0	0
## 10		165.100	54.48774	54.0	1
## 11		154.305	49.89512	47.0	0
## 12		151.130	41.22017	66.0	1
## 13		144.780	36.03221	73.0	0
## 14		149.900	47.70000	20.0	0
## 15		150.495	33.84930	65.3	0

# DATA

- For now, focus on data for individuals under age 15.

```
htm <- Howell1$height/100
bmi <- Howell1$weight/(htm^2)
y_male <- bmi[Howell1$age<15 & Howell1$male==1]
y_female <- bmi[Howell1$age<15 & Howell1$male==0]
n_m <- length(y_male)
n_f <- length(y_female)

n_f
```

```
## [1] 84
```

```
n_m
```

```
## [1] 77
```

```
summary(y_male)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    12.07   13.87   14.63   14.84   15.53   18.22
```

```
summary(y_female)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##     9.815  13.559  14.305  14.585  15.712  18.741
```

# CLASSICAL INFERENCE

- No significant difference in group means.

```
t.test(y_male,y_female)
```

```
##  
##      Welch Two Sample t-test  
##  
## data:  y_male and y_female  
## t = 1.1204, df = 157.87, p-value = 0.2643  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  -0.1947946  0.7054729  
## sample estimates:  
## mean of x mean of y  
##  14.84037  14.58503
```

# SIMPLE WEIGHTED ESTIMATOR

- One parameterization that can reflect some relationship between  $\theta_m$  and  $\theta_f$  is

$$\begin{aligned} y_{i,\text{male}} &\stackrel{iid}{\sim} \mathcal{N}(\mu + \delta, \sigma^2); \quad i = 1, \dots, n_m; \\ y_{i,\text{female}} &\stackrel{iid}{\sim} \mathcal{N}(\mu - \delta, \sigma^2); \quad i = 1, \dots, n_f. \end{aligned}$$

where

- $\theta_m = \mu + \delta$  and  $\theta_f = \mu - \delta$ ,
- $\mu = \frac{\theta_m + \theta_f}{2}$  is the pooled average, and
- $\delta = \frac{\theta_m - \theta_f}{2}$  is half of the population difference in means.

# SIMPLE WEIGHTED ESTIMATOR

- Convenient prior:
  - $\pi(\mu, \delta, \sigma^2) = \pi(\mu) \cdot \pi(\delta) \cdot \pi(\sigma^2)$ , where
    - $\pi(\mu) = \mathcal{N}(\mu_0, \gamma_0^2)$ ,
    - $\pi(\delta) = \mathcal{N}(\delta_0, \tau_0^2)$ , and
    - $\pi(\sigma^2) = \mathcal{IG}(\frac{\nu_0}{2}, \frac{\nu_0 \sigma_0^2}{2})$ .
- We will set the hyper-parameters as:
  - $\mu_0 = 15, \gamma_0 = 5$ ,
  - $\delta_0 = 0, \tau_0 = 3$ ,
  - $\nu_0 = 1, \sigma_0 = 5$ .
- Do these values seem reasonable to you?

# SIMPLE WEIGHTED ESTIMATOR

- Note that we can rewrite

$$\begin{aligned} y_{i,\text{male}} &\stackrel{iid}{\sim} \mathcal{N}(\mu + \delta, \sigma^2); \quad i = 1, \dots, n_m; \\ y_{i,\text{female}} &\stackrel{iid}{\sim} \mathcal{N}(\mu - \delta, \sigma^2); \quad i = 1, \dots, n_f \end{aligned}$$

as

$$\begin{aligned} (y_{i,\text{male}} - \delta) &\stackrel{iid}{\sim} \mathcal{N}(\mu, \sigma^2); \quad i = 1, \dots, n_m; \\ (y_{i,\text{female}} + \delta) &\stackrel{iid}{\sim} \mathcal{N}(\mu, \sigma^2); \quad i = 1, \dots, n_f \end{aligned}$$

or

$$\begin{aligned} (y_{i,\text{male}} - \mu) &\stackrel{iid}{\sim} \mathcal{N}(\delta, \sigma^2); \quad i = 1, \dots, n_m; \\ (-1)(y_{i,\text{female}} - \mu) &\stackrel{iid}{\sim} \mathcal{N}(\delta, \sigma^2); \quad i = 1, \dots, n_f. \end{aligned}$$

as needed, so we can leverage past results for the full conditionals.

# FULL CONDITIONALS

- For the full conditionals we will derive today, we will take advantage of previous results from the regular univariate normal model.
- Recall that if we assume

$$y_i \sim \mathcal{N}(\mu, \sigma^2), \quad i = 1, \dots, n,$$

and set our priors to be

$$\begin{aligned}\pi(\mu) &= \mathcal{N}(\mu_0, \gamma_0^2) . \\ \pi(\sigma^2) &= \mathcal{IG}\left(\frac{\nu_0}{2}, \frac{\nu_0 \sigma_0^2}{2}\right),\end{aligned}$$

then we have

$$\pi(\mu, \sigma^2 | Y) \propto \left\{ \prod_{i=1}^n p(y_i | \mu, \sigma^2) \right\} \cdot \pi(\mu) \cdot \pi(\sigma^2)$$



# FULL CONDITIONALS

- We have

$$\pi(\mu|\sigma^2, Y) = \mathcal{N}(\mu_n, \gamma_n^2).$$

where

$$\gamma_n^2 = \frac{1}{\frac{n}{\sigma^2} + \frac{1}{\gamma_0^2}}; \quad \mu_n = \gamma_n^2 \left[ \frac{n}{\sigma^2} \bar{y} + \frac{1}{\gamma_0^2} \mu_0 \right],$$

- and

$$\pi(\sigma^2|\mu, Y) = \mathcal{IG}\left(\frac{\nu_n}{2}, \frac{\nu_n \sigma_n^2}{2}\right),$$

where

$$\nu_n = \nu_0 + n; \quad \sigma_n^2 = \frac{1}{\nu_n} \left[ \nu_0 \sigma_0^2 + \sum_{i=1}^n (y_i - \mu)^2 \right].$$

# FULL CONDITIONALS

- With  $\pi(\mu) = \mathcal{N}(\mu_0, \gamma_0^2)$ , we have

$$\mu|Y, \delta, \sigma^2 \sim \mathcal{N}(\mu_n, \gamma_n^2), \quad \text{where}$$

$$\gamma_n^2 = \frac{1}{\frac{1}{\gamma_0^2} + \frac{n_m + n_f}{\sigma^2}}$$

$$\mu_n = \gamma_n^2 \left[ \frac{\mu_0}{\gamma_0^2} + \frac{n_m \overline{(y_{i,male} - \delta)} + n_f \overline{(y_{i,female} + \delta)}}{\sigma^2} \right].$$

where

- $\overline{(y_{i,male} - \delta)} = \frac{1}{n_m} \sum_{i=1}^{n_m} (y_{i,male} - \delta)$ , and
- $\overline{(y_{i,female} + \delta)} = \frac{1}{n_f} \sum_{i=1}^{n_f} (y_{i,female} + \delta)$ .

# FULL CONDITIONALS

- With  $\pi(\delta) = \mathcal{N}(\delta_0, \tau_0^2)$ , we have

$$\delta|Y, \mu, \sigma^2 \sim \mathcal{N}(\delta_n, \tau_n^2), \quad \text{where}$$

$$\tau_n^2 = \frac{1}{\frac{1}{\tau_0^2} + \frac{n_m + n_f}{\sigma^2}}$$

$$\delta_n = \tau_n^2 \left[ \frac{\delta_0}{\tau_0^2} + \frac{n_m \overline{(y_{i,male} - \mu)} + (-1)n_f \overline{(y_{i,female} + \mu)}}{\sigma^2} \right].$$

where

- $\overline{(y_{i,male} - \mu)} = \frac{1}{n_m} \sum_{i=1}^{n_m} (y_{i,male} - \mu)$ , and
- $\overline{(y_{i,female} - \mu)} = \frac{1}{n_f} \sum_{i=1}^{n_f} (y_{i,female} - \mu)$ .

# FULL CONDITIONALS

- With  $\pi(\sigma^2) = \mathcal{IG}(\frac{\nu_0}{2}, \frac{\nu_0 \sigma_0^2}{2})$ , we have

$$\sigma^2 | Y, \mu, \delta \sim \mathcal{IG}(\frac{\nu_n}{2}, \frac{\nu_n \sigma_n^2}{2}), \text{ where}$$

$$\nu_n = \nu_0 + n_m + n_f$$

$$\sigma_n^2 = \frac{1}{\nu_n} \left[ \nu_0 \sigma_0^2 + \sum_{i=1}^{n_m} (y_{i,male} - [\mu + \delta])^2 + \sum_{i=1}^{n_f} (y_{i,female} - [\mu - \delta])^2 \right].$$

# APPLICATION TO DATA

```
#priors
mu0 <- 15; gamma02 <- 5^2
delta0 <- 0; tau02 <- 3^2
nu0 <- 1; sigma02 <- 5^2

#starting values
mu <- (mean(y_male) + mean(y_female))/2
delta <- (mean(y_male) - mean(y_female))/2
#no need for starting values for sigma_squared, we can sample it first

MU <- DELTA <- SIGMA2 <- NULL
```

# APPLICATION TO DATA

```
#set seed
set.seed(1234)

#set number of iterations and burn-in
n_iter <- 10000; burn_in <- 0.2*n_iter

##Gibbs sampler
for (s in 1:(n_iter+burn_in)) {
  #update sigma2
  sigma2 <- 1/rgamma(1,(nu0 + n_m + n_f)/2,
                    (nu0*sigma02 + sum((y_male-mu-delta)^2) + sum((y_female-mu+delta)^2))/2)

  #update mu
  gamma2n <- 1/(1/gamma02 + (n_m + n_f)/sigma2)
  mun <- gamma2n*(mu0/gamma02 + sum(y_male-delta)/sigma2 + sum(y_female+delta)/sigma2)
  mu <- rnorm(1,mun,sqrt(gamma2n))

  #update delta
  tau2n <- 1/(1/tau02 + (n_m+n_f)/sigma2)
  deltan <- tau2n*(delta0/tau02 + sum(y_male-mu)/sigma2 - sum(y_female-mu)/sigma2)
  delta <- rnorm(1,deltan,sqrt(tau2n))

  #save parameter values
  MU <- c(MU,mu); DELTA <- c(DELTA,delta); SIGMA2 <- c(SIGMA2,sigma2)
}
```

# POSTERIOR SUMMARIES

```
#library(coda)
MU.mcmc <- mcmc(MU,start=1)
summary(MU.mcmc)
```

```
##
## Iterations = 1:12000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 12000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean           SD      Naïve SE Time-series SE
##      14.712517      0.118765      0.001084      0.001089
##
## 2. Quantiles for each variable:
##
##  2.5%   25%   50%   75%  97.5%
## 14.48 14.63 14.71 14.79 14.95
```

```
(mean(y_male) + mean(y_female))/2 #compare to data
```

```
## [1] 14.7127
```

# POSTERIOR SUMMARIES

```
DELTA.mcmc <- mcmc(DELTA,start=1)
summary(DELTA.mcmc)
```

```
##
## Iterations = 1:12000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 12000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean           SD       Naive SE Time-series SE
##      0.127657      0.119522      0.001091      0.001091
##
## 2. Quantiles for each variable:
##
##      2.5%      25%      50%      75%      97.5%
## -0.10691  0.04791  0.12743  0.20796  0.36407
```

```
summary((2*DELTA)) #rescale as difference in group means
```

```
##      Min.  1st Qu.   Median     Mean  3rd Qu.    Max.
## -0.63464  0.09582  0.25487  0.25531  0.41592  1.23660
```

```
mean(y_male) - mean(y_female) #compare to data
```

```
## [1] 0.2553392
```



# POSTERIOR SUMMARIES


```
SIGMA2.mcmc <- mcmc(SIGMA2,start=1)
summary(SIGMA2.mcmc)
```

```
##
## Iterations = 1:12000
## Thinning interval = 1
## Number of chains = 1
## Sample size per chain = 12000
##
## 1. Empirical mean and standard deviation for each variable,
##    plus standard error of the mean:
##
##           Mean           SD      Naive SE Time-series SE
##      2.287927      0.257689      0.002352      0.002352
##
## 2. Quantiles for each variable:
##
##  2.5%   25%   50%   75%  97.5%
## 1.833 2.107 2.272 2.455 2.841
```


# DIAGNOSTICS

```
plot(MU.mcmc)
```

**Trace of var1**




**Density of var1**



N = 12000 Bandwidth = 0.01924

# DIAGNOSTICS


```
autocorr.plot(MU.mcmc)
```




# DIAGNOSTICS

```
plot(DELTA.mcmc)
```

**Trace of var1**




**Density of var1**



# DIAGNOSTICS


```
autocorr.plot(DELTA.mcmc)
```




# DIAGNOSTICS

```
plot(SIGMA2.mcmc)
```

**Trace of var1**




**Density of var1**



N = 12000 Bandwidth = 0.04174

# DIAGNOSTICS

```
autocorr.plot(SIGMA2.mcmc)
```



# APPLICATION TO DATA

- Posterior probability that boys have larger average BMI than girls is 0.86!
- Posterior medians and 95% credible intervals for the group means are actually quite similar to the unpooled (gender specific) intervals from classified inference.

```
#mean for boys  
quantile((MU+DELTA),probs=c(0.025,0.5,0.975))
```

```
##      2.5%      50%      97.5%  
## 14.50255 14.84146 15.17925
```

```
#mean for girls  
quantile((MU-DELTA),probs=c(0.025,0.5,0.975))
```

```
##      2.5%      50%      97.5%  
## 14.26848 14.58276 14.90761
```

```
#posterior probability girls have larger BMI than boys  
mean(DELTA > 0)
```

```
## [1] 0.8571667
```



# APPLICATION TO DATA

- Let's look at a different sub-population. For older individuals  $> 75$ , we only have 8 male and 4 female.

```
y_male <- bmi[Howell1$age > 75 & Howell1$male==1]  
y_female <- bmi[Howell1$age > 75 & Howell1$male==0]  
n_m <- length(y_male)  
n_f <- length(y_female)  
n_m
```

```
## [1] 8
```

```
n_f
```

```
## [1] 4
```

# APPLICATION TO DATA

- A 95% confidence interval for the difference between genders in BMI (estimated as 0.24) is  $(-4.20, 4.68)$ .

```
mean(y_male) - mean(y_female)
```

```
## [1] 0.2408966
```

```
t.test(y_male,y_female)
```

```
##  
##      Welch Two Sample t-test  
##  
## data:  y_male and y_female  
## t = 0.13801, df = 5.1869, p-value = 0.8954  
## alternative hypothesis: true difference in means is not equal to 0  
## 95 percent confidence interval:  
##  -4.197948  4.679741  
## sample estimates:  
## mean of x mean of y  
##  18.06751  17.82662
```

# APPLICATION TO DATA

- Let's apply the shrinkage model with these priors:
  - $\mu_0 = 18, \gamma_0 = 5,$
  - $\delta_0 = 0, \tau_0 = 3,$
  - $\nu_0 = 1, \sigma_0 = 5.$
- Using the shrinkage model, the posterior mean is 0.25 with 95% CI (-3.45, 3.88).

```
mean((DELTA*2))
```

```
## [1] 0.2493733
```

```
quantile((DELTA*2), probs=c(0.025, 0.5, 0.975))
```

```
##          2.5%          50%          97.5%  
## -3.4466931  0.2758598  3.8762543
```

- Our precision has been improved by borrowing of information across the groups. Of course the prior is important here given the sample sizes.

# COMPARING MULTIPLE GROUPS

- Suppose we wish to investigate the mean (and distribution) of test scores for students at  $J$  different high schools.
- In each school  $j$ , where  $j = 1, \dots, J$ , suppose we test a random sample of  $n_j$  students.
- Let  $y_{ij}$  be the test score for the  $i$ th student in school  $j$ , with  $i = 1, \dots, n_j$ , with

$$y_{ij} | \theta_j, \sigma_j^2 \sim \mathcal{N}(\theta_j, \sigma_j^2)$$

where for each school  $j$ ,  $\theta_j$  is the school-wide average test score, and  $\sigma_j^2$  is the school-wide variance of individual test scores.

- This is what we did for the the Pygmalion study, job training data and the science classroom exercise on homework 3.

# SCHOOL TESTING EXAMPLE

- Classical inference for each school can be based on large sample 95% CI:  $\bar{y}_j \pm 1.96 \sqrt{s_j^2/n_j}$ , where  $\bar{y}_j$  is the sample average in school  $j$ , and  $s_j^2$  is the sample variance in school  $j$ .
- Clearly, we can overfit the data within schools, for example, what if we only have 4 students from one of the schools?  $\bar{y}_j$  can be a good estimate if  $n_j$  is large but it may be poor if  $n_j$  is small.
- **Option II**: alternatively, we might believe that  $\theta_j = \mu$  for all  $j$ ; that is, all schools have the same mean. This is the assumption (null hypothesis) in ANOVA models for example. We can also set  $\sigma_j^2 = \sigma^2$  for all  $J$ .
- Option I ignores that the  $\theta_j$ 's should be reasonably similar, whereas option II ignores any differences between them.
- It would be nice to find a compromise! Borrowing information across, and shrinking our estimate towards a **grand mean** could be very useful here.

# SCHOOL TESTING EXAMPLE

- For the Pygmalion study and job training data, we focused using priors that are independent between the groups.
- For example, in the conjugate case, we would have

$$\begin{aligned}\pi(\theta_j | \sigma_j^2) &= \mathcal{N}\left(\mu_0, \frac{\sigma_j^2}{\kappa_0}\right) \\ \pi(\sigma_j^2) &= \mathcal{IG}\left(\frac{\nu_0}{2}, \frac{\nu_0 \sigma_0^2}{2}\right)\end{aligned}$$

for some hyperparameters (constants),  $\mu_0$ ,  $\kappa_0$ ,  $\nu_0$ , and  $\sigma_0^2$ .

- In the semi-conjugate case,

$$\begin{aligned}\pi(\theta_j) &= \mathcal{N}(\mu_0, \sigma_0^2) \\ \pi(\sigma_j^2) &= \mathcal{IG}\left(\frac{\nu_0}{2}, \frac{\nu_0 \gamma_0^2}{2}\right)\end{aligned}$$

for some hyperparameters (constants),  $\mu_0$ ,  $\sigma_0^2$ ,  $\nu_0$ , and  $\gamma_0^2$ .

# HIERARCHICAL NORMAL MODEL

- Instead, we can assume that the  $\theta_j$ 's are drawn from a distribution based on the following: conceive of the schools themselves as being a random sample from all possible schools.
- For now, assume the variance is constant across schools. The hierarchical normal model assumes normal sampling models both within and between groups:

$$\begin{aligned}y_{ij}|\theta_j, \sigma^2 &\sim \mathcal{N}(\theta_j, \sigma^2); \quad i = 1, \dots, n_j \\ \theta_j|\mu, \tau^2 &\sim \mathcal{N}(\mu, \tau^2); \quad j = 1, \dots, J,\end{aligned}$$

which gives us an extra level in the prior on the means, which leads to sharing of information across the groups in estimating the group-specific means.

- We have an extra variance parameter  $\tau^2$ . Comparing  $\tau^2$  to  $\sigma^2$  tells us how much of the variation in  $Y$  is due to within-group versus between-group variation.

# HIERARCHICAL NORMAL MODEL

- Standard semi-conjugate priors are given by

$$\begin{aligned}\pi(\mu) &= \mathcal{N}(\mu_0, \gamma_0^2) \\ \pi(\sigma^2) &= \mathcal{IG}\left(\frac{\nu_0}{2}, \frac{\nu_0 \sigma_0^2}{2}\right) \\ \pi(\tau^2) &= \mathcal{IG}\left(\frac{\eta_0}{2}, \frac{\eta_0 \tau_0^2}{2}\right).\end{aligned}$$

with

- $\mu_0$ : best guess of average of school averages
- $\gamma_0^2$ : set based on plausible ranges of values of  $\mu$
- $\tau_0^2$ : best guess of variance of school averages
- $\eta_0$ : set based on how tight prior for  $\tau^2$  is around  $\tau_0^2$
- $\sigma_0^2$ : best guess of variance of individual test scores around respective school means
- $\nu_0$ : set based on how tight prior for  $\sigma^2$  is around  $\sigma_0^2$ .



# EXCHANGEABILITY

- This model relies heavily on exchangeability across units at each level.
- For example, we assume the schools are a random sample from the population of all schools, and the students within schools are a random sample of all the students in each school.
- This is not always completely true.
- Note: we can allow the variance to vary across schools if desired (and we will soon in fact).

# EXCHANGEABILITY

- Turns out that **conditional exchangeability** would be enough if we control for relevant variables in our modeling.
- For example, the schools in Chapel Hill/Carrboro are not entirely exchangeable.
- For example, Phoenix Academy is for students on long-term out-of-school suspension or who need to make up work due to extended absences (e.g., pregnancy), and Memorial Hospital School is for children battling serious illnesses.
- However, if we condition on school type (public, charter, private, special services, home), the schools may then be exchangeable.

# POSTERIOR INFERENCE

- Recall the model is

$$\begin{aligned}y_{ij}|\theta_j, \sigma^2 &\sim \mathcal{N}(\theta_j, \sigma^2); \quad i = 1, \dots, n_j \\ \theta_j|\mu, \tau^2 &\sim \mathcal{N}(\mu, \tau^2); \quad j = 1, \dots, J,\end{aligned}$$

- Under our prior specification, we can factor the posterior as follows:

$$\begin{aligned}\pi(\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2|Y) &\propto p(y|\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2) \\ &\quad \times p(\theta_1, \dots, \theta_J|\mu, \sigma^2, \tau^2) \\ &\quad \times \pi(\mu, \sigma^2, \tau^2) \\ &= p(y|\theta_1, \dots, \theta_J, \sigma^2) \\ &\quad \times p(\theta_1, \dots, \theta_J|\mu, \tau^2) \\ &\quad \times \pi(\mu) \cdot \pi(\sigma^2) \cdot \pi(\tau^2) \\ &= \left\{ \prod_{j=1}^J \prod_{i=1}^{n_j} p(y_{ij}|\theta_j, \sigma^2) \right\} \\ &\quad \times \left\{ \prod_{j=1}^J p(\theta_j|\mu, \tau^2) \right\} \\ &\quad \times \pi(\mu) \cdot \pi(\sigma^2) \cdot \pi(\tau^2)\end{aligned}$$

# FULL CONDITIONAL FOR GRAND MEAN

- The full conditional distribution of  $\mu$  is proportional to the part of the joint posterior  $\pi(\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2 | Y)$  that involves  $\mu$ .
- That is,

$$\pi(\mu | \theta_1, \dots, \theta_J, \sigma^2, \tau^2, Y) \propto \left\{ \prod_{j=1}^J p(\theta_j | \mu, \tau^2) \right\} \cdot \pi(\mu).$$

- This looks like the full conditional distribution from the one-sample normal case, so you can show that

$$\pi(\mu | \theta_1, \dots, \theta_J, \sigma^2, \tau^2, Y) = \mathcal{N}(\mu_n, \gamma_n^2) \quad \text{where}$$

$$\gamma_n^2 = \frac{1}{\frac{J}{\tau^2} + \frac{1}{\gamma_0^2}}; \quad \mu_n = \gamma_n^2 \left[ \frac{J}{\tau^2} \bar{\theta} + \frac{1}{\gamma_0^2} \mu_0 \right]$$

$$\text{and } \bar{\theta} = \frac{1}{J} \sum_{j=1}^J \theta_j.$$

# FULL CONDITIONALS FOR GROUP MEANS

- Similarly, the full conditional distribution of each  $\theta_j$  is proportional to the part of the joint posterior  $\pi(\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2 | Y)$  that involves  $\theta_j$ .
- That is,

$$\pi(\theta_j | \mu, \sigma^2, \tau^2, Y) \propto \left\{ \prod_{i=1}^{n_j} p(y_{ij} | \theta_j, \sigma^2) \right\} \cdot p(\theta_j | \mu, \tau^2)$$

- Those terms include a normal for  $\theta_j$  multiplied by a product of normals in which  $\theta_j$  is the mean, again mirroring the one-sample case, so you can show that

$$\pi(\theta_j | \mu, \sigma^2, \tau^2, Y) = \mathcal{N}(\theta_j^*, \nu_j^*) \quad \text{where}$$

$$\nu_j^* = \frac{1}{\frac{n_j}{\sigma^2} + \frac{1}{\tau^2}}; \quad \theta_j^* = \nu_j^* \left[ \frac{n_j}{\sigma^2} \bar{y}_j + \frac{1}{\tau^2} \mu \right]$$

# FULL CONDITIONALS FOR GROUP MEANS

- Our estimate for each  $\theta_j$  is a weighted average of  $\bar{y}_j$  and  $\mu$ , ensuring that we are borrowing information across all levels through  $\mu$  and  $\tau^2$ .
- The weights for the weighted average is determined by relative precisions from the data and from the second level model.
- The groups with smaller  $n_j$  have estimated  $\theta_j^*$  closer to  $\mu$  than schools with larger  $n_j$ .
- Thus, degree of shrinkage of  $\theta_j$  depends on ratio of within-group to between-group variances.

# FULL CONDITIONALS FOR ACROSS-GROUP VARIANCE

- The full conditional distribution of  $\tau^2$  is proportional to the part of the joint posterior  $\pi(\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2 | Y)$  that involves  $\tau^2$ .
- That is,

$$\pi(\tau^2 | \theta_1, \dots, \theta_J, \mu, \sigma^2, Y) \propto \left\{ \prod_{j=1}^J p(\theta_j | \mu, \tau^2) \right\} \cdot \pi(\tau^2)$$

- As in the case for  $\mu$ , this looks like the one-sample normal problem, and our full conditional posterior is

$$\pi(\tau^2 | \theta_1, \dots, \theta_J, \mu, \sigma^2, Y) = \mathcal{IG} \left( \frac{\eta_n}{2}, \frac{\eta_n \tau_n^2}{2} \right) \quad \text{where}$$

$$\eta_n = \eta_0 + J; \quad \tau_n^2 = \frac{1}{\eta_n} \left[ \eta_0 \tau_0^2 + \sum_{j=1}^J (\theta_j - \mu)^2 \right].$$

# FULL CONDITIONALS FOR WITHIN-GROUP VARIANCE

- Finally, the full conditional distribution of  $\sigma^2$  is proportional to the part of the joint posterior  $\pi(\theta_1, \dots, \theta_J, \mu, \sigma^2, \tau^2 | Y)$  that involves  $\sigma^2$ .
- That is,

$$\pi(\sigma^2 | \theta_1, \dots, \theta_J, \mu, \tau^2, Y) \propto \left\{ \prod_{j=1}^J \prod_{i=1}^{n_j} p(y_{ij} | \theta_j, \sigma^2) \right\} \cdot \pi(\sigma^2)$$

- We can again take advantage of the one-sample normal problem, so that our full conditional posterior is

$$\pi(\sigma^2 | \theta_1, \dots, \theta_J, \mu, \tau^2, Y) = \mathcal{IG} \left( \frac{\nu_n}{2}, \frac{\nu_n \sigma_n^2}{2} \right) \quad \text{where}$$

$$\nu_n = \nu_0 + \sum_{j=1}^J n_j; \quad \sigma_n^2 = \frac{1}{\nu_n} \left[ \nu_0 \sigma_0^2 + \sum_{j=1}^J \sum_{i=1}^{n_j} (y_{ij} - \theta_j)^2 \right].$$